

The HBR Editorial Standard

A Complete Writing System Extracted from Harvard Business Review

For use as the quality filter in AI-assisted book production. This standard defines what the output must sound like, how it must be structured, what evidence it requires, and what it must never do.

PART I: VOICE AND PROSE STYLE

The Sound

HBR sounds like a very smart person explaining something important to another very smart person over a serious dinner. It is neither the journal nor the airport bookstore. The vocabulary assumes graduate-level education but avoids academic obscurantism.

Sentence Rhythm

The dominant pattern: ****short declarative lead, one or two longer elaborating sentences, then another short punch.**** PUNCH -- expand -- expand -- PUNCH.

- Average sentence length: 15-22 words, with wide range tolerated
- Long sentences (30+) are fine if syntactically clean
- Compound-complex sentences with multiple subordinate clauses are not
- Subject-Verb-Object backbone, main clause front-loaded
- Semicolons: sparingly. Em dashes: more common. Colons: freely for setups

Paragraph Shape

- Body: 3-6 sentences, 80-120 words
- Openings: shorter (2-4 sentences) to pull in fast
- Single-sentence paragraphs: for emphasis, no more than 1-2 per article
- Never exceed 150 words without breaking

Vocabulary

- Target: well-read MBA or senior executive
- Terms like "salient," "exogenous," "asymmetric" appear without definition
- Specialized academic terms get immediate practical context
- Prefer concrete Anglo-Saxon verbs ("cut," "build," "shift") over Latinate abstractions ("facilitate," "utilize," "implement")

Person and Voice

- ****Practitioner pieces****: first person freely ("When I became CEO...")
- ****Research pieces****: "In our research" for process, third person for findings
- ****Second person "you"****: strategically in prescriptive sections
- ****Imperative mood****: common in how-to sections ("Start by mapping your stakeholders.")
- ****Passive voice****: avoided. "Our research found" not "It was found"
- ****Pattern****: narrative opening (first/third) -> analysis (third) -> prescription (second/imperative)

Tone Coordinates

- ****Authoritative but accessible**** (65-70 on a 0-100 academic-to-conversational scale)
- ****Confident but not arrogant**** -- "here is what we know" not "here is what I, a very important person, think"
- ****Urgent but not breathless**** -- something is at stake, and here is what to do
- ****Controlled warmth**** -- occasional dry wit, never jokes with punchlines, no sarcasm

What HBR Does NOT Sound Like

- A textbook (too lively, too opinionated)
- A blog post (too structured, too evidence-based)
- A newspaper (too analytical, too prescriptive)
- A TED talk (too detailed, less focused on inspiration)

- McKinsey Quarterly (warmer, more narrative)

Hedging vs. Declarative

HBR leans ****strongly declarative****.

- One hedge per claim maximum. Double-hedging is edited out
- Preferred hedges: "often," "typically," "in many cases"
- "Could" and "might" mostly for forward-looking statements
- Prescriptive passages are declarative/imperative: "Leaders need X" not "Leaders might want to consider X"
- Qualifiers go in separate bounded sentences, not weakening the main claim

****Confidence gradient across an article:****

- Opening claim: bold, declarative
- Evidence section: slightly qualified ("our research suggests")
- Prescriptive section: returns to declarative/imperative
- Conclusion: bold again, often bolder than the opening

Rhetorical Questions

- 2-3 per article maximum
- Always answered immediately -- within same paragraph or next sentence
- Never stacked
- Appear at paragraph beginnings, not endings
- Used as section transitions: "So how should leaders respond?"

Transitions

1. ****Explicit pivot****: "But strategy is only part of the equation." / "That's the diagnosis. Here's the prescription."
2. ****Subheadings****: every 400-700 words. Short, declarative, not cute
3. ****Bridge paragraphs****: 1-2 sentences linking sections
4. ****Framework introductions****: "We've identified four types of..."
5. ****Adversative pivots****: "here's what most people think" -> "but/however" -> "here's what is actually true"

Avoid: "Furthermore," "moreover," "additionally" -- these signal listing, not arguing.

Attribution

- Research: natural language, not citations. "Research by Amy Edmondson of Harvard Business School shows that..."
- Credentials: front-loaded and brief, then surname only
- Data: attributed in same sentence. "According to Gartner, 85% of..."
- Self-attribution: direct. "In our study of 147 companies..."
- Quotes: minimal setup. "As one CEO told us, '[quote].'"
- ****No footnotes, no parenthetical citations, no ibid.****

PART II: STRUCTURE

The Arc

Every significant HBR piece follows this proportional structure:

Phase	Weight	Function
Hook	5-10%	Establish the problem. Make the reader feel it
Diagnosis	25-35%	Why it exists. What conventional wisdom gets wrong
Framework	20-30%	The new model, tool, or way of thinking
Evidence	15-25%	Case studies and data proving the framework
Prescription	10-20%	What to do about it. Steps, guidelines, implications
Wider Lens	3-5%	Zoom out. Larger implications. Forward-looking

The Four Opening Moves

1. **The Executive Anecdote**: Named individual facing a specific moment of decision. Most common by far
2. **The Counterintuitive Claim**: Bold statement challenging conventional wisdom
3. **The Trend with Stakes**: Widespread phenomenon that is bigger than you think
4. **The Research Hook**: Surprising finding from a large study

Never open with: dictionary definitions, grand philosophical statements, "In today's rapidly changing business environment," or the author's credentials.

Section Rhythm (for a ~5,000 word feature)

- **Opening** (600-900 words): hook, problem, thesis preview. No subheading
- **Diagnostic sections** (2-3, 500-800 words each): build the case
- **Framework section** (800-1200 words): the named concept. Often with exhibit
- **Application sections** (1-2, 400-700 words each): framework in action
- **Conclusion** (300-500 words): shortest section. Zoom out

Frameworks and Models

HBR is where business concepts get named. The atomic unit is the **named, reusable intellectual tool**:

- Numbered steps (Kotter's 8 Steps)
- Named components (Goleman's 5 components of EI)
- 2x2 matrices (BCG Growth-Share)
- Diagnostic canvases (Blue Ocean's ERRC grid)

A good framework is **non-obvious and discriminating** -- it tells you what to do differently, not what everyone already agrees matters.

Closings

- Zoom out to bigger picture
- Restate core insight in sharpest form
- Forward-looking note
- **Do NOT**: summarize point by point, introduce new evidence, end with generic optimism, end with a question

The "Big Idea" Formula

Nearly every successful HBR article follows this deep structure:

*Here is something you think you understand. You are partially wrong. Here is the more nuanced truth. Here is what it means for how you lead/manage/decide. Here are the specific things to do differently.**

For Book-Length Work

- Chapters: 3,000-5,000 words each (3,000-4,000 average)
- Each chapter advances one major idea within the larger thesis
- Structure resembles expanded feature articles: opening scene, conceptual argument, evidence, application
- More case studies per chapter than per article (5-10 vs 3-6)
- Deeper implementation detail
- Each chapter should stand alone as an argument while building the book's arc

PART III: EVIDENCE AND RIGOR

The Evidence Hierarchy

Tier	Type	How Used
1	Original research (peer-reviewed, proprietary datasets)	Lead with finding, not methodology. Cite conclusion and sample size, not p-value
2	Named case studies (real companies, real people, real outcomes)	Specific, verifiable, 3+ from

different industries to establish pattern |

| 3 | Expert testimony (named executives, practitioners) | Direct experience as evidence, grounded in specifics |

| 4 | Quantitative data points | Statistics as anchors, not as the argument. Longitudinal preferred |

| 5 | Historical parallels | Used sparingly, with explicit acknowledgment of where the analogy breaks down |

What Does NOT Count as Evidence

- Unsourced assertions ("Everyone knows...")
- Anecdotes without names or specifics
- Circular references to other thought leadership
- Theoretical frameworks without empirical grounding
- Pure logic without real-world validation

The Data Sandwich

1. Open with provocative quantitative claim (stakes)
2. Humanize with specific story
3. Return to data showing the story isn't an outlier
4. Close with actionable qualitative insight

Numbers create credibility; stories create understanding. Never let either stand alone.

Citing Research Without Sounding Academic

- Lead with finding, not researcher: "People feel losses roughly twice as powerfully as equivalent gains"
- Name researchers as characters: "Amy Edmondson spent two decades studying..."
- Translate methodology: "researchers randomly assigned managers to two groups"
- Omit what practitioners don't need (statistical methods, literature reviews)
- Name institutions when they add credibility

What Needs a Source vs. What Can Be Asserted

****Must have evidence:****

- Any specific number
- Any causal claim
- Any claim about what "most" do
- Any claim contradicting conventional wisdom (higher bar, not lower)

****Can be asserted:****

- Widely accepted principles
- Logical deductions from established premises
- Author's own framework (marked as proposed model)
- Observations clearly labeled as such

Handling Counterarguments

- Acknowledge, don't dwell: one paragraph
- Steelman selectively -- only real business objections a smart executive would raise
- Use boundary conditions instead of hedging: "This works when X. In organizations that Y, a different strategy may be needed."
- The "To be sure..." construction: acknowledges complexity, then returns to main argument

Predictions and Forward-Looking Claims

- Must be falsifiable
- Specify what they depend on
- Acknowledge alternative scenarios
- Scenario framing preferred over point predictions
- Confidence signaled through structural choices, not just language

PART IV: WHAT HBR IS AND IS NOT (Positioning)

HBR's Unique Position

Axis	HBR's Position	
-----	-----	
Academic rigor <->	Practitioner access	Midpoint
Timely <->	Timeless	Heavily timeless
Descriptive <->	Prescriptive	Heavily prescriptive
Individual voice <->	Institutional voice	Fully institutional

vs. The Economist

HBR chose the manager as reader; The Economist chose the informed citizen. HBR asks "What should I do Monday morning?" The Economist asks "What is actually happening?"

vs. McKinsey Quarterly

HBR chose independence; McKinsey chose access. HBR can challenge any worldview. McKinsey can never truly challenge the worldview its clients pay for.

vs. MIT Sloan Management Review

HBR chose the practitioner; SMR chose the practitioner-scholar. HBR sacrifices nuance for clarity. SMR sacrifices accessibility for accuracy.

vs. The Atlantic / New Yorker

HBR chose utility; literary journalism chose understanding. HBR builds toward recommendation. Literary journalism builds toward illumination.

vs. Fast Company / Wired

HBR chose durability; Fast Company chose currency. HBR is building a library. Fast Company is building a feed.

vs. Substack / Medium

HBR chose the institution; independent writers chose the individual. Different trust architectures producing different knowledge.

The Reader's Expectations

1. A clear, nameable idea they can apply
2. Evidence, not just opinion
3. Respectful peer-to-peer address
4. Actionability -- do something different after reading
5. Authority earned through research or practice
6. Optimistic pragmatism -- "This is hard, and here is how"
7. Compression -- density with no wasted paragraphs

What Feels Off-Brand

- Polemics, anger, or cynicism
- Pure narrative with no framework
- Technical jargon without translation
- Pessimism without prescription
- Personal essay (unless illuminating a management principle)
- Humor as primary mode
- Ideological advocacy
- Speculation without grounding
- Celebrity worship

PART V: THE BLACKLIST

Tier 1: Kill On Sight (Empirically Verified AI Tells)

delve, meticulous, commendable, intricate, showcasing, underscores, tapestry, multifaceted, nuanced (as compliment)

Tier 2: High-Frequency LLM Words (Use Sparingly If At All)

landscape (as metaphor), navigate/navigating (as metaphor), leverage/leveraging, comprehensive, pivotal, crucial, notable/notably, enhancing, realm, interplay

Tier 3: AI Filler Connectors

Furthermore, Moreover, Indeed, Certainly, Undoubtedly, It is worth noting that, It is important to note that, Importantly, Significantly, Remarkably, Interestingly

Banned Opening Lines

- "In today's rapidly changing world..."
- "In an era of unprecedented change..."
- "As we stand at the crossroads of..."
- "In the ever-evolving world of..."
- "Now more than ever..."

Banned Phrases

- "game-changing," "groundbreaking," "paradigm shift," "cutting-edge," "best-in-class"
- "best practices," "key takeaways," "actionable insights," "unlock value," "move the needle"
- "thought leadership," "synergy," "ecosystem," "solution" (replacing product)
- "reimagine," "lean into," "purpose-driven," "unlock potential"
- "The future belongs to those who..."
- "Only time will tell"

Silicon Valley Vocabulary to Avoid

- "disrupt" (unless citing Christensen's actual theory)
- "10x" as adjective, "moonshot," "first principles thinking," "founder mode"
- "ship fast and break things," "zero to one," "blitzscaling"

Structural Anti-Patterns

****Openings that fail:**** dictionary definitions, historical panoramas, statistics bombardment, rhetorical question salvos, vague anecdotes, throat-clearing paragraphs

****Middles that fail:**** the accordion (same idea expanding/contracting), the catalog (examples without analysis), pivot avoidance, invisible logic, padding

****Endings that fail:**** restatement, unearned uplift, punting ("only time will tell"), calls to "reimagine," false certainty, summary paragraphs

AI Structural Tells

- Symmetrical lists (every section has exactly 3 points)
- The hedging swamp ("while there are challenges, there are also opportunities")
- The missing "I" (avoiding first person)
- Uniform paragraph length
- Unnaturally positive tone throughout
- Over-signposting ("First we'll examine... Then we'll explore...")

PART VI: THE TESTS

The Five HBR Editorial Tests

1. ****Expertise****: Demonstrated, not claimed. Early establishment of credentials
2. ****Evidence****: Research, data, case studies validating the argument
3. ****Originality****: New idea or genuinely new perspective. Most common rejection reason
4. ****Usefulness****: Applicable in business today. The "Monday morning test"
5. ****Good Writing****: Clear, persuasive, a pleasure to read

The Seven Pitch Questions (Every Piece Must Answer)

1. What is your central message?
2. What's important, useful, or counterintuitive about it?
3. What makes it novel?
4. Why do senior leaders need this now?
5. What is your authority to write this?
6. What research and examples support your claims?
7. What relevant experience will you draw upon?

The 20-Question Kill Screen

1. Could ChatGPT produce this paragraph from a simple prompt? If yes, rewrite
2. Does this contain a Tier 1 banned word? If yes, replace
3. Does this say something the reader already knows? If yes, cut
4. Would a busy executive read past the first line? If no, rewrite
5. Does this make a claim without evidence? If yes, add evidence or remove
6. Could this appear on any company's corporate blog? If yes, make it specific
7. Does this sound like a TED talk? If yes, add mechanism and specificity
8. Does this sound like an academic paper? If yes, add examples and cut jargon
9. Is there a "both sides" hedge avoiding a position? If yes, commit
10. Does this use a dying metaphor or cliché? If yes, say it fresh or plain
11. Does this assume all readers are tech founders? If yes, broaden
12. Can a sentence be cut without losing meaning? If yes, cut
13. Passive voice without good reason? Make it active
14. First person missing where expertise is claimed? Add "I"
15. Does the section end stronger than it begins? If not, restructure
16. Would the reader know what to do differently? If not, add specificity
17. Does the anecdote have enough detail to be credible? If vague, detail or cut
18. Does this engage the strongest counterargument? If not, find and address it
19. Would a domain expert learn something new? If not, go deeper
20. Does this respect reader intelligence without excluding them? Fix if it patronizes or obscures

The Authenticity Test

The standard is not "was AI involved?" but "does this contain genuine human insight that AI could not have produced alone?"

For every paragraph: Does this contain something that could only come from a person who has actually experienced, built, or deeply studied this subject? If a person with no domain expertise could produce substantially similar text by prompting an LLM, it fails.

The Real Expertise Markers (vs. Performed Expertise)

Real operators have:

- Scar tissue (stories of failures, pilots that cratered)
- Altitude flexibility (macro trends AND granular operational details)
- Specific failure narratives (not "we learned a lot")
- Evolved perspectives (mental models that changed over time)
- Accountability (skin in the game through implementation)

Performers have:

- All-wins narratives
- Comfort only at 30,000 feet
- Generic failure statements
- Rapidly accrued expertise
- Consistency suggesting brand protection over learning

PART VII: THE THEORY-PRACTICE RATIO

The typical HBR ratio:

- ****~40% diagnostic/analytical****: what's happening and why
- ****~35% framework/model****: how to think about it

- ****~25% prescriptive/actionable****: what to do about it

HBR almost never prescribes without first earning the right through rigorous diagnosis. This ratio distinguishes it from McKinsey (more prescriptive), academic journals (more diagnostic), and trade publications (almost purely prescriptive).

Orwell's Rules (Still the Foundation)

1. Never use a metaphor, simile, or figure of speech you are used to seeing in print
2. Never use a long word where a short one will do
3. If it is possible to cut a word out, always cut it out
4. Never use the passive where you can use the active
5. Never use a foreign phrase, scientific word, or jargon word if you can think of an everyday English equivalent
6. Break any of these rules sooner than say anything outright barbarous

****The core principle:**** "Good prose is like a windowpane." The reader should see through the writing to the ideas. When the writing itself becomes visible -- through jargon, hedging, pretension, or cliché -- it has failed.

This standard was compiled from HBR contributor guidelines, editor interviews (Adi Ignatius, Sarah Green Carmichael, Amy Bernstein), published HBR author advice, the 2024 Stanford/Tubingen study of AI-generated text markers, Orwell's "Politics and the English Language," Paul Graham's writing advice, and structural analysis of canonical HBR articles (Porter, Christensen, Goleman, Kotter, Kim & Mauborgne).